2nd CAPSTONE PROJECT STATUS CHECKING FORM

Project Title:	CS-57 [AI / Software Engineering] Human-AI Data Annotation, Provenance, and Dataset Evaluation Platform
Project Client:	Hunter Xu
Group Number:	CS-57
Date of client meeting:	23-September-2026

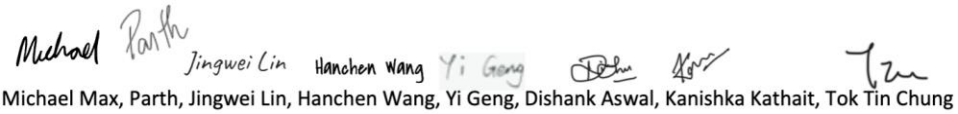
Status Highlight	**Summary of report & feedback**
Overall Project Status	
Progress and Achievements	• **Basic workflow works end to end** 1. Invited into an organization → Admin & Task owner create project → create task → add items into task (done) 2. Annotators perform data annotation → Reviewers review, then accept/reject→ send item to dispute if conflicts happen between reviewers (in-progress) 3. Expert (Arbitrator) resolves the dispute — a basic ruling works; the client's revised version (accept, return or reject) is in progress → item finalised (done) 4. Admin & Task owner export reviewed and finalised items (done); the full provenance history with each exported item (to-do) • **AI-Provider integration complete** — typed provider interface with mock and live interchangeable, batch jobs, per-item job status, background worker, retry that separates transient from permanent failures • **Trustworthy Foundations complete** — legacy review mutations retired behind 410 Gone, cross-project and cross-org write protection, environment and seed-data repair • **Annotation & Human Judgement** — draft ownership, atomic draft submission, custom task and label definitions with schema upload • **Organization management backend done** (members, invitations, roles); project review-policy settings mostly done; task status changes (activate, pause, complete) in progress this sprint • **Migrated to PostgreSQL 18** while keeping SQLite working — local development stays fast, CI tests the production engine, enabling concurrent data processing in multi-user and batch scenarios • **452 backend tests** pass (4 skipped) and 205 frontend tests across 27 files on the main branch as of 23 September; every push runs the backend tests against both SQLite and PostgreSQL, with PostgreSQL set to a non-UTC timezone to catch date bugs • **Governance rules enforced on the server,** not only hidden in the UI — for example, only the author can edit their draft, users cannot write into another organisation's projects, reviewers cannot approve their own annotations, and a finalised item rejects further edits; the same self-approval check for disputed items is being added this sprint

	Client's feedback: Your understanding to the project now is deeper, while still on the correct and critical path, I especially appreciate that you are working on the task queue and dispute and review workflow now. The change from SQLite to PostgreSQL makes sense to me, while I would suggest to have a sentence justifying the choice in your final report.
Key Issues	
Obstacles & Risks	• **Raised at the first checking — inherited defects:** 30 defects were catalogued in the inherited code; 20 are fixed or have a fix merged, including 3 of the 6 critical ones. As the client directed, the rest are fixed when a feature being built depends on them, not in a separate clean-up phase. • **Raised at the first checking — no review workflow:** resolved. Every change now goes through a pull request reviewed by another team member (25 merged since 6 September), and automated tests run on every push; the one gap is GitHub branch protection, below. • **Evaluation and provenance not started:** there is no code yet for the evaluation harness (automated re-runs of the whole workflow), the casebook of 40–60 test cases, or the provenance log (a per-item record of who did what, when, and under which guideline). These are heavily weighted deliverables, with half the semester left. • **Release depends on provenance:** the release export planned for week 10 must list each item's history and refuse items whose history is incomplete, but the provenance log is currently scheduled for week 11. • **Design questions settled with the client as work reaches them:** some expected behaviours can only be confirmed once a feature is being built — for example, what happens after an expert rules on a dispute, who may reopen a finalised item, and how an uploaded dataset is split into items. Waiting for an answer can hold a user story, and an answer can reveal work that was not in the backlog; this week's answers added three new stories. This tests the team's agile planning: stories have to be re-scoped, and the remaining weeks re-planned without losing momentum.
	Client's feedback: Good reflections on these, especially what you pointed out are quite important and valuable, another piece of evidence that you grasped the main line of the project. Regarding the evaluation and harness, I would suggest to have the use cases in mind and note them down while you are developing the relevant features, a backlog of cases would then benefit the build of the case book and tests.
Deviation from Timeline	• **Against the plan presented at the first checking:** user stories and the pull-request workflow (week 5) are done; the core annotation backend (weeks 7–8) is done apart from task status changes, in progress this week; review, roles and dispute work (weeks 9–10) started early and is in progress; the end-to-end run on the provided datasets (week 11) needs JSONL import, brought forward to this sprint. The separate schema and contract review (week 6) was replaced by fixing each mismatch inside the feature that needs it. • **Scope since the first checking:** the first checking agreed to deliver the core workflow first and treat further functionality as extension work, so provenance, releases and evaluation were not in that plan. The brief and the client's answers since then confirm they are required, so they are now scheduled from the mid-semester break to week 11.

	• **Core platform:** on track — the workflow works end to end up to review, while dispute handling, expert adjudication and export with provenance are incomplete; AI pre-annotation is complete apart from the batch progress screen (under review), and the PostgreSQL migration landed without stalling feature work • **Evaluation and provenance:** behind — both unstarted with the back half of semester remaining. This is a deviation in sequencing rather than total effort: the team front-loaded the platform and deferred the assessment-weighted deliverables. For evaluation the brief asked for the opposite ordering • **Review and adjudication:** in progress — draft ownership, the self-approval check, sending work back to the annotator and locking finalised items are merged; this sprint adds a work queue where annotators and reviewers pick up available items, an independent second review for a sample of items, and a separate record of each expert ruling. • A few week-7 items carried over into week 8 — the work-queue API, member-management screens, project-policy settings and the self-approval check for disputes; the other deviations are in work not yet started, not in work delivered late
	Client's feedback: Looks good to me, I reckon the project is still align with the overall schedule, keep building.
Next steps	
Milestones	**From the plan presented at the first checking** 1. **Week 5** User stories and the pull-request review workflow — Achieved 2. **Week 6** Schema and contract validation — Replaced: mismatches are fixed inside each feature 3. **Weeks 7–8** Core annotation backend (task setup, AI annotation, drafts) — Achieved, plus AI provider integration and the PostgreSQL migration; task status changes in progress 4. **Weeks 9–10** Human review, role-based governance and dispute resolution — In progress, started early in week 7 5. **Week 11** End-to-end run on the provided datasets — In progress: JSONL import brought forward to week 8 6. **Week 12** Hardening, documentation and demo preparation — Planned 7. **Week 13** Final presentation and report — Planned **Added since the first checking (brief and client answers)** 8. **Mid-semester break** Evaluation harness with repeatable runs and regression reporting — Not started 9. **Weeks 10–11** Casebook of 40–60 adversarial and realistic cases — Not started; 40 cases by week 10 10. **Weeks 9–11** Provenance log, versioned sources and immutable releases with manifests — Not started; the provenance log is currently in week 11 and should move earlier
	Client's feedback: From what I can sense, some of the modules require deep and detailed design (e.g. schema, workflow, authentication, provenance). If you find time is tight and the task is hard, we can have a working prototype for these modules at first and expand the details in an iterative manner, such that every time since then we would have a running workflow to test the features end-to-end. We don't have to have perfect design and implementation at the start of each task.

How to address obstacles & risks mentioned above in the 'Key Issues'	• **Evaluation first in the mid-semester break:** the break is used to build the harness, and every bug found while scenario-testing becomes a new case; target 40 cases by week 10. • **Provenance log before the release export:** move the provenance log ahead of the week-10 release work, so the export reads real item history instead of having it added afterwards. • **Decisions recorded as they arrive:** each client decision is written up as an architecture decision record (a short note of the decision, the alternatives and the reasons), and the affected backlog items and tickets are updated the same day; further questions go to the client as GitHub issues. • **Repository controls:** every pull request runs the backend tests on both databases and, by team rule, is reviewed by someone other than its author; the server blocks write across projects, and blocking reads across projects follows after the core workflow, if time allows.
	Client's feedback: Sounds good, let's make it together.

All group member signatures (either handwritten or digital signatures):



Michael Max, Parth, Jingwei Lin, Hanchen Wang, Yi Geng, Dishank Aswal, Kanishka Kathait, Tok Tin Chung

Client signatures (either handwritten or digital signatures):



Weekly Client Meeting Minutes:

Time: No client meeting held. Written progress update issued Sep. 6, 2026

Venue: Email

Meeting Minute Taker: Michael

Attendances: Hunter Xu (client); Michael (author)

Apologies: N/A

Main Discussions / What Happened

- No client meeting was scheduled this week. With the scope signed off at the check-in on Aug. 26, the week was given to internal work - breaking the scope into a structured backlog, finalising the Group Proposal and allocating role responsibilities. A meeting before that work was complete would have had little new to report.
- Michael issued a written progress update to the client covering the backlog structure, role allocation, the Jira board and the defect handling approach.
- Michael proposed a meeting on Wednesday Sep. 9 as an early checkpoint on the first sprint, allowing the client to review the backlog and delivery plan and raise any concerns before implementation is far advanced.

Key Takeaways

- The agreed scope is now decomposed into 59 stories across 11 epics, with inherited defects linked to the stories they affect.
- Implementation begins next week; next client aimed for Sep. 9.

What's Next

- Michael is responsible for delivering the W5 update via email to the client.
- Michael is responsible for confirming the Wednesday Sep. 9 meeting with the client.

We acknowledge Mr. Hunter Xu for assistance in preparing this template.

Weekly Client Meeting Minutes:

Time: Sep. 9, 2026, the 1st meeting

Venue: Zoom meeting

Meeting Minute Taker: Michael, Hanchen

Attendances: Everyone

Apologies: N/A

Main Discussions / What Happened

- Michael explained why no client meeting was scheduled the previous week: with the inherited codebase in poor condition, the time was better spent understanding it and identifying tasks than meeting before there was anything new to report.
- Hanchen presented the team's artefacts to the client - the Jira board for task allocation, the shared spreadsheet for contribution logs, and the user story dashboard demonstrating coverage of the deliverables for later weeks.
- Yi raised a list of questions on the technical details of the data annotation workflow.
- The client cautioned that Jira can become difficult to manage, as tasks can be split into subtasks indefinitely, and advised the team not to spend disproportionate effort keeping the board tidy.

Key Takeaways

- The weekly client meeting is confirmed for every Wednesday at 7pm for the remainder of the project.
- The client directed the team to prioritise delivering features in GitHub over documenting work in Jira. Tracking should remain sufficient for coordination rather than exhaustive.

What's Next

- All team members are responsible for delivering their allocated Jira tasks ahead of next week's meeting.
- The team is responsible for presenting the contribution log and a short demonstration of progress at next week's client meeting.

We acknowledge Mr. Hunter Xu for assistance in preparing this template.

Weekly Client Meeting Minutes:

Time: Sep. 16, 2026, the 1st meeting

Venue: Zoom meeting

Meeting Minute Taker: Michael, Yi

Attendances:

Apologies: Tim Chung

Main Discussions / What Happened

- Hanchen presented the team's progress on the project, summarising the sprint content, the GitHub pull requests and contribution logs, and the Jira dashboard, with a live demonstration.
- Yi explained the integration design of the MVP repository and the results. The client was satisfied with the AI model selection function, but noted that the currently enumerated subtasks may not adapt well to a future increase in the number of task types.
- Hanchen and Yi raised a list of questions on the technical details of the data annotation workflow and task management.
- Given that questions and answers raised during the meeting may not be captured clearly, the client suggested the team create an online Google Doc to record questions and answers systematically.

Key Takeaways

- The client clarified that a cross-marking function is required: multiple annotators must be able to work on the same item and create their own drafts simultaneously.
- The client suggested that where a user does not have permission for a function, the button should be greyed out rather than displaying a notification.

What's Next

- All team members are responsible for delivering their allocated Jira tasks ahead of next week's meeting.
- The team is responsible for presenting the contribution log and a short demonstration of progress at next week's client meeting.

We acknowledge Mr. Hunter Xu for assistance in preparing this template.

Weekly Client Meeting Minutes:

Time: Sep. 23, 2026, the 1st meeting

Venue: Zoom meeting

Meeting Minute Taker: Michael, Yi

Attendances:

Apologies: Kanishka

Main Discussions / What Happened

- Reported the migration to PostgreSQL 18 with SQLite retained for local development, and the current test position: 452 backend tests passing across both databases and 205 frontend tests, with PostgreSQL run in a non-UTC timezone to catch date handling bugs.
- The team explained that the majority of this sprint's work was backend, so while substantial progress was made there was comparatively little to demonstrate visually.
- Reported the AI annotation pipeline/queue as complete - typed provider interface with mock and live interchangeable, batch jobs, per-item job status, a background worker, and retry logic separating transient from permanent failures.
- Worked through the Q&A document with the client, resolving outstanding design questions directly rather than deferring them to implementation.
- Discussed the mid-semester break. The team confirmed it will continue at reduced capacity rather than stopping, with the evaluation harness as the priority for that period.

Key Takeaways

- The Q&A document is working as intended: design and archetectual questions are resolved as work reaches them rather than being deferred. This week's answers added three new stories to the backlog.

What's Next

- All team members are responsible for delivering their allocated W8 tickets by 30 September.
- The team is responsible for starting the evaluation harness over the break, targeting 40 casebook cases by W10.
- Client provides feedback and sign-off for Project Status Checking 2

We acknowledge Mr. Hunter Xu for assistance in preparing this template.